

[Best Read on GitHub!](#)

This Submission in a nutshell

(time to complete: 2hrs 10mins)

"Stand on the shoulders of giants" (particularly when cognitively depleted, compute deprived, and time constrained)

[Live development of this submission.](#) Recommend you only watch the first 3 mins and last 3 mins (lost in the middle like 😊)

Task 1

Response

(RESPONSE GENERATED BY kapa.ai, If this does not resolve your problem please dont hesitate to reply again with more details and a member of the weaviate team will be happy to help)

Yes, Weaviate does support multimodal search. It allows you to store data objects and vector embeddings from your favorite ML-models. You can use additional modalities such as images, audio, or video as the query, and retrieve corresponding, compatible objects. Some modules, such as multi2vec-clip and multi2vec-bind allow you to search across modalities. For example, you can search for images using a text query, or search for text using an image query [\[source\]](#). Weaviate also supports image, or multimodal data and performs the same type of "near image" search across a vector space [\[source\]](#). To get started with Weaviate, you can check out the [Quickstart tutorial](#), or build amazing apps with a free trial of [Weaviate Cloud \(WCD\)](#) [\[source\]](#). For a more practical approach, you can follow the steps in the Weaviate Academy's guide on working with multimodal data. It includes units on setting up a local Weaviate instance, populating the database, performing searches, and using large language models to augment and transform retrieved data [\[source\]](#). Here is an example of how to create a Multimodal Collection in Weaviate:

```
client.collections.create(  
    name="Animals",  
    vectorizer_config=wvc.config.Configure.Vectorizer.multi2vec_bind(  
        audio_fields=["audio"],  
        image_fields=["image"],  
        video_fields=["video"],  
    )  
)
```

)

[\[source\]](#) And here is an example of how to perform an image search:

```
response = animals.query.near_image(  
    near_image=toBase64("./test/test-cat.jpg"),  
    return_properties=['name', 'path', 'mediaType'],  
    limit=3  
)
```

[\[source\]](#) I hope this helps you get started with implementing multimodal RAG into your product using Weaviate!

Tickets

- Create a Reddit app to interface kapa.ai into Weaviate's supreddit by listening to posts (at least). If sufficiently robust extending the kapa.ai app to Reddit should be trivial.

Task 2

Weaviate: 1, Waldo: 0 🏆 Our vector database just beat Jonathan's kid at finding Waldo! Learn form how this fun experiment showcases the power of AI in complex image analysis using @weaviate_io.

<https://weaviate.io/blog/find-waldo-with-weaviate>

#Weaviate #AIphoto #GenerativeAI



Expanding

Create an interactive experience for this use case where we deploy the app that does this and we tap into the X api to listen to replies from users where they can reply to the post with an "Where's Waldo" landscape image and the X bot replies with the relative location of Waldo (although this is pretty time consuming I would imagine it would result in a unforgettable experience for X users).

Ready to challenge Weaviate's Waldo-finding skills? 🕵️🔍 Reply to this tweet with your toughest 'Where's Waldo?' image, and our AI will try to spot him! #WaldoAIChallenge #WeaviateVision

X user responds with:

Here

[IMAGE]

Waldo X Bot replies with:

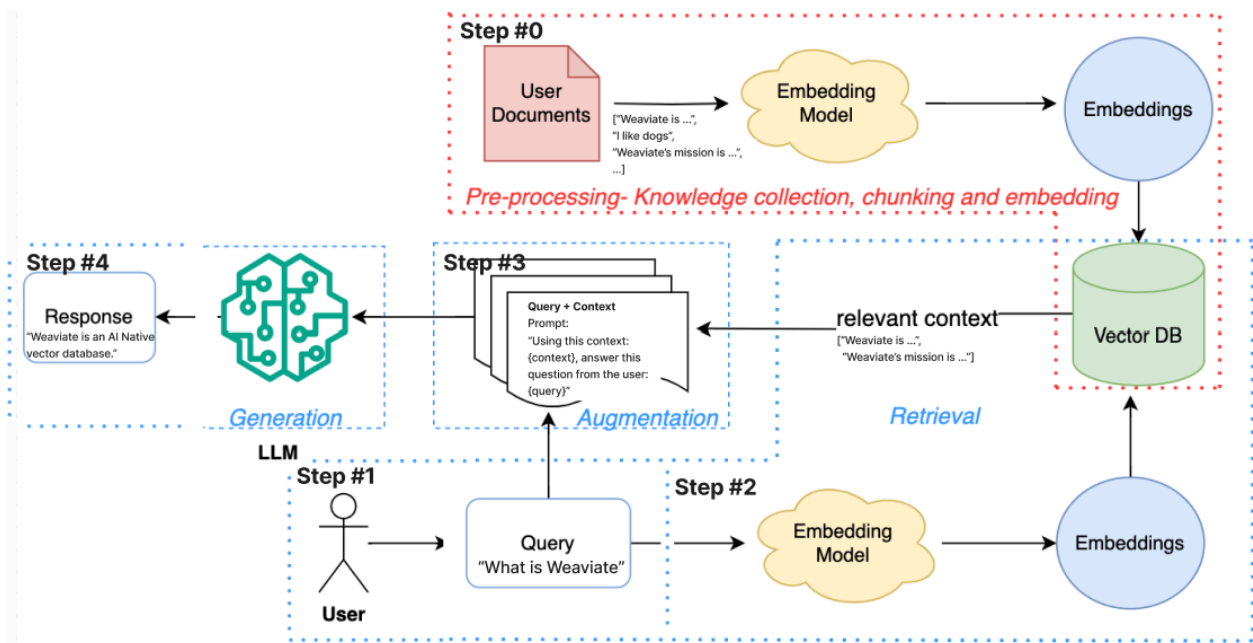
Waldo is found at normalized X coordinate .9 and normalized Y coordinate .2

[IMAGE_OF_WALDO_FROM_USER_IMAGE]

Tickets

- Create the aforementioned app

Task 3



Base diagram src: <https://mindfulmatrix.substack.com/p/build-a-simple-llm-application-with>